


Beam search over the target vocabulary



Scores are cumulative log-probabilities — every one is negative, so a longer hypothesis is penalised purely for being longer. Dividing by $\text{length}^{0.6}$ before the final comparison corrects that; without it beam search systematically truncates.